

Technical Report for the IMU Odometry Challenge

Trajectory-Context Inertial Velocity Estimation with Metric-Shaped Training

Kaggle team name: **Hack2Publish**

Md. Hamid Hosen¹ and Esfer Sami²

Abstract—We describe a single unified network (one shared set of weights for car, quadruped, drone and handheld) that predicts the mean body-frame velocity of every 1 s window from raw 6-axis IMU alone. The core idea is to stop treating windows in isolation: the model reads 16 s chunks of a trajectory (past *and* future), predicts a dense 20 Hz velocity for the whole chunk, and overlapping chunks are averaged at inference. The trunk is a strided convolutional stem augmented with deterministic strap-down features, eight dilated depthwise temporal-convolution blocks and a three-layer transformer encoder (3.9 M parameters). Training is shaped after the metric: platforms are sampled uniformly and trajectories uniformly within a platform, and an integrated-error loss targets ATE_{20} directly. Augmentation is physically exact (random sensor re-mount rotations applied to both the IMU and the target velocity, time dilation using the training-time gravity vector, bias/scale/noise). All model selection was done on the held-out val split with the organizers’ scorer; the final model was refit on train+val for 160 epochs with EMA and SWA. The submitted model scores 0.219 on the full 89-sequence test set (official scoring service; macro AVE 0.180 m/s, ATE_{20} 0.563 m; car/dog/human AVE 0.05–0.07 m/s, drone 0.52 m/s) and 0.28609 on the public leaderboard (rank 9 of 108 at close, from an initial 0.553 of ours), with a held-out val score of 0.198. We also report two data findings that bound what any IMU-only model can achieve on the drone split, and a measured public-leaderboard noise of ± 0.01 –0.02 for identical recipes.

I. INTRODUCTION

Learning an inertial odometer that works across embodiments is hard for three reasons. (i) The observability of velocity from a 6-axis IMU is fundamentally limited: gravity, sensor bias and mount orientation are entangled with the signal, and slow motions (a hovering drone, a person standing) produce almost no measurable acceleration. (ii) The four platforms differ by an order of magnitude in speed, vibration and dynamics, and the test set is platform-blind, so one network must infer the embodiment from the signal itself. (iii) The ground truth is not homogeneous: it comes from different capture systems with different conventions and timing.

Our solution was motivated by a simple observation in the training data: the per-window velocity is strongly autocorrelated (lag-1 correlation 0.85 car, 0.78 dog, 0.56 human), and an oracle that merely averages the two neighbouring windows’ true velocities already halves the all-zero error. Since test trajectories are given whole, a model is allowed to use context; a 1 s window cannot estimate gyroscope bias,

low-frequency tilt or the embodiment reliably, but 16 s can. The key components are therefore (1) a trajectory-context architecture with dense output and overlap averaging, (2) a training procedure shaped after the macro-/per-trajectory averaging of the TartanIMU score, and (3) physically exact augmentation. Our contributions beyond the model are two data findings (Sec. IV-D) that explain most of the drone difficulty and a measurement of leaderboard noise that affected how we selected our final entries.

II. METHOD

A. Overall Architecture

Fig. 1 shows the pipeline. Input is a raw 200 Hz chunk of $T=16$ consecutive 1 s windows (3200×6 , $[a_x, a_y, a_z, \omega_x, \omega_y, \omega_z]$, gravity retained), divided channel-wise by $(3, 3, 3, 0.4, 0.4, 0.4)$ with no mean subtraction. A strided convolutional stem (kernels 9/2 and 11/5) produces 20 Hz tokens; deterministic strap-down features computed on the same chunk (gyro-integrated relative orientation via a parallel prefix product of per-token rotations, the acceleration rotated into the chunk-start frame with the chunk mean taken as a gravity estimate, and its integral, all expressed back in the current body frame; computed under both gyro- z sign hypotheses, see Sec. IV-D) are projected and added to the tokens. Eight residual depthwise-dilated temporal-convolution blocks (dilations 1,2,4,8,16,32,1,2; receptive field ≈ 13 s, GroupNorm) are followed by a 3-layer transformer encoder (width 192, 4 heads, learned positions) over the whole chunk and a per-token linear head that outputs dense 20 Hz body-frame velocity, averaged to one vector per window. An attention-pooled descriptor feeds an auxiliary 4-way platform classifier used only as a training signal. Nothing in the network is conditioned on a platform label; the embodiment is inferred implicitly, which the auxiliary head merely encourages. Design choices aimed at cross-platform generalization: no per-channel mean normalization (keeps gravity as a tilt cue and makes rotation augmentation exact), bidirectional context (bias and tilt are estimated from the whole chunk), and platform-balanced sampling.

B. Data Processing

Windows are used exactly as defined ($k \cdot 200$ to $(k+1) \cdot 200$); trajectories are kept whole in memory and training samples are random 16-window chunks (edge-padded and loss-masked when a trajectory is shorter). Dense 20 Hz

¹Md. Hamid Hosen, team leader. hamidhosen8444@gmail.com

²Esfer Sami, team member.




Fig. 1. Overview of the method. One network, one weight set; the auxiliary platform head is discarded at inference.

targets are the 10-sample means of `vel.body`. Augmentation, all on device and per chunk: **sensor re-mount rotation** (uniform random axis, angle $U(0, 15^\circ)$; $U(0, 45^\circ)$ for drone chunks) applied identically to accelerometer, gyroscope and target velocity, i.e. an exact re-mounting; **time dilation** by $s \sim \log U(1/1.3, 1.3)$: the chunk is resampled, $v \rightarrow sv$, $\omega \rightarrow s\omega$ and the dynamic acceleration $\rightarrow s^2$ while the gravity component (from the training quaternion) is kept; accelerometer scale $1 + \mathcal{N}(0, 0.02)$, bias $\mathcal{N}(0, 0.15 \text{ m/s}^2)$, gyroscope scale $1 + \mathcal{N}(0, 0.02)$, bias $\mathcal{N}(0, 0.02 \text{ rad/s})$, white noise $\sigma = 0.05 / 0.004$. No filtering, resampling or calibration is applied at inference.

C. Training Data

Only the official challenge data was used (Table I). Model selection used the released `train/val` split (395/80 trajectories); the final model was refit on `train+val` with the schedule fixed by the `val` experiments. No external data, synthetic data, pretrained checkpoints or pseudo-labels. Test information was never used: test files were read only by the inference script.

TABLE I
DATASETS USED FOR MODEL DEVELOPMENT AND TRAINING.

Dataset	Type	Usage
Challenge <code>train</code>	Official	Training of all selection runs
Challenge <code>val</code>	Official	Model/recipe selection (never trained on in selection runs); folded into the final refit

D. Training and Optimization

Loss: $L = L_{\text{win}} + 0.5 L_{\text{dense}} + 0.2 L_{\text{drift}} + 0.05 L_{\text{plat}}$, where L_{win} is a vector Huber ($\beta = 0.25 \text{ m/s}$) on the per-window mean velocity, L_{dense} the same on the 20 Hz output, $L_{\text{drift}} = \frac{1}{T} \sum_t \|\sum_{k \leq t} (\dot{v}_k - v_k)\| / \sqrt{t}$ an integrated body-frame error mirroring ATE_{20} 's sensitivity to correlated bias, and L_{plat} the auxiliary cross-entropy. AdamW ($\beta = 0.9 / 0.99$,

weight decay 0.02), OneCycle (peak 1.5×10^{-3} , 8 % warm-up), batch 64 chunks, 250 steps/epoch, gradient clipping 2.0, EMA of the weights (0.998) and a uniform average of the EMA weights over the last 20 epochs (SWA). Sampling: platform uniform, then trajectory uniform within the platform. Selection runs: 60 epochs; final refit: 160 epochs, single stage, no curriculum, no fine-tuning, no platform-specific specialization. The final checkpoint is the SWA weights at the end of the fixed schedule – no early stopping is possible because `val` is inside its training set. Hardware: Apple M5 laptop GPU for all selection runs (1.9–3 min/epoch); the ranked checkpoint was trained by a Kaggle script kernel on a Kaggle GPU in 3.6 h.

E. Inference

For each test trajectory, 16-window chunks slide with a stride of 2 windows; each chunk is run once; per-window outputs are averaged across overlapping chunks with a raised-cosine weight (+0.05 floor) so a window is trusted most from the chunk in which it is central. Trajectories shorter than 16 windows are edge-padded and the padding discarded. Trajectories are processed independently. No TTA, no ensembling, no scaling, smoothing or clipping, no test-time adaptation. Inference uses past and future context of the same trajectory (offline).

F. Model Complexity and Runtime

Table II. The method is offline (bidirectional); a causal variant would need a different design.

TABLE II
MODEL COMPLEXITY AND RUNTIME STATISTICS.

Property	Value
Number of parameters	3.9 M (second entry: 1.8 M)
Training hardware	Kaggle GPU (ranked run); Apple M5 laptop GPU (selection runs)
Training time	3.6 h (160 epochs $\times$ 250 steps $\times$ 64 chunks)
Inference hardware	any single GPU (<2 GB VRAM) or CPU
Inference speed (Hz)	≈ 500 windows/s on a laptop GPU (whole test set <1 min); CPU: tens of minutes
Peak memory	<2 GB GPU; ≈ 3 GB RAM

III. DEVELOPMENT EXPERIMENTS

All numbers below are `val` scores (organizers' scorer, 80 held-out trajectories) of models trained on `train` only, unless stated. Fig. 2 shows the curves.

Context is the largest gain. A per-window 1D ResNet in the style of the public starter kernel reaches ≈ 0.30 ; our first 16-s-context model (width 128, 30 epochs) reached **0.2247**. A 32-s context with a wider trunk was *worse* at equal epochs (0.2293), so 16 s was kept.

Longer schedules. The same recipe for 60 instead of 30 epochs: 0.2233 $\rightarrow$ **0.2030** (drone-B AVE 0.328 $\rightarrow$ 0.271, car 0.129 $\rightarrow$ 0.117, human 0.085 $\rightarrow$ 0.073). On the public leaderboard the `train+val` refits went 0.3625 (30 ep) $\rightarrow$ 0.3466 (60) $\rightarrow$ 0.2975 (100) $\rightarrow$ 0.2870 (160); 240 epochs gave no further gain (0.2897).

Metric-shaped sampling. Sampling trajectories uniformly within a platform instead of proportionally to length: 0.2030→**0.1982**, with the largest effect on the short, hard drone trajectories – the metric averages per trajectory, so short trajectories count as much as long ones.

Model scale. Width 192 with a 3-layer transformer instead of 128/2: 0.2030→0.1987 on val, but the corresponding train+val refits scored 0.2893 and 0.2861–0.2876 publicly versus 0.2870 for the narrow model, i.e. within leaderboard noise (Sec. VI-B). Scale helped every platform slightly and hurt none.

Drone-targeted augmentation. Rotations up to 45° for drone chunks, ×4 oversampling of the racing-drone source and time dilation: 0.2247→0.2233 overall but *no change* on the racing-drone trajectories (AVE 0.674→0.676). Increasing the drone sampling share to 40% (0.2013) and vibration-amplitude augmentation (0.2065) were both worse than uniform sampling (0.1982).

Things that did not work (cost in parentheses): raw strap-down features alone gave no measurable gain – 16 s integration drifts by 5–20 m/s under uncompensated gyro bias, so the network’s implicit short-horizon integration is already better (2 h); a learned per-chunk IMU↔ground-truth *time-shift head*, supervised by measured offsets, predicted the offsets of unseen recordings well (−66/−48 ms predicted vs −40/−70 measured) but did not improve the score (2 h); rotation TTA (0.2246 vs 0.2247), 3-window smoothing (ATE worse on every platform), denser overlap (identical), a bidirectional GRU context module (same accuracy, 3× slower on our hardware). Prediction averaging of 2–4 runs improved val to 0.2171 and public to 0.354 (from 0.363) but is not a single weight set and was not used for any ranked entry.




Fig. 2. Validation score and drone AVE during training for the main selection runs (train only, scored on the held-out val split every 2–4 epochs).

IV. EXPERIMENTS

A. Official Challenge Performance

Tables III–V were produced by the official scoring service for the CSV of our ranked Kaggle submission 56402554, md5 = **fd12d4f25d7b52d44aaa4e4a14534378** (scored under the team’s pre-rename Kaggle name “Md. Hamid Hosen”; the team was renamed Hack2Publish on 21 Sept. and the service’s whitelist had not yet refreshed). Full-test score **0.21903**; public leaderboard 0.28609 (rank 9 of 108 at close). The second selected entry (56332266, md5 97a6e86dba7dfc19929e050bef5857b0) scores 0.22189 on the same service (AVE 0.181, ATE₂₀ 0.581). The drone term carries 64% of our score (Table IV): AVE

is 0.05–0.07 m/s on car, human and dog but 0.52 m/s on drone, and ATE₂₀ 0.36–0.49 m versus 1.04 m.

TABLE III

OVERALL CHALLENGE PERFORMANCE (OFFICIAL SCORING SERVICE, FULL 89-SEQUENCE TEST SET). THE TARTANIMU SCORE IS 0.6 AVE/0.7356 + 0.4 ATE₂₀/3.1160, DIMENSIONLESS, *lower is better* (ALL-ZERO = 1.000). **RTE DOES NOT ENTER THE RANKING**; IT IS DEFINED IN THE CAPTION OF TABLE V. BASELINE NUMBERS ARE THE ORGANIZERS’ PUBLISHED FULL-TEST VALUES.

Method	Score ↓	ATE ₂₀ ↓	AVE ↓	RTE ↓
TartanIMU Baseline	0.538	1.261	0.461	–
Ours (Hack2Publish)	0.2190	0.563	0.180	0.868

TABLE IV

PERFORMANCE BY PLATFORM (OFFICIAL SCORING SERVICE). RTE IS DEFINED IN THE CAPTION OF TABLE V.

Platform	ATE ₂₀ ↓	AVE ↓	RTE ↓
Wheeled / Car	0.369	0.071	0.323
Handheld / Human	0.490	0.053	0.270
Legged / Quadruped	0.358	0.074	0.353
Aerial / Drone	1.037	0.522	2.528
Average	0.563	0.180	0.868

B. Per-Sequence Results

The error is concentrated, not spread (Table V): a minority of drone sequences – those of the racing-drone source (short, 3–7 m/s, two IMUs per flight, Sec. IV-D) – have AVE several times that of any other sequence and account for most of the drone term, which itself is 64% of the score. Within the other platforms the per-sequence errors are homogeneous (car/dog/human AVE 0.07–0.13 m/s in val); the slow hovering drone sequences have an AVE equal to the standard deviation of their own ground-truth velocity, i.e. the model predicts the mean and cannot resolve sub-0.2 m/s hover drift from the IMU.

C. Ablation Studies

Section III is our ablation record; in decreasing order of effect on val: trajectory context (−0.08), schedule length (−0.02), uniform trajectory sampling (−0.005), model scale (−0.004), drone augmentation (−0.001), physics features / time-shift head / TTA / smoothing (0 or negative).

D. Generalization Analysis

Val (disjoint trajectories) tracks the test set closely: val 0.198 vs 0.219 on the full test set, while the public split is much harder (0.286); the organizers’ baseline shows the same pattern (0.538 full test vs 0.637 public). The difference is explained by the drone split. Fig. 3 shows the per-trajectory val error of our recipe: car, dog and human errors are 0.07–0.13 m/s regardless of speed; drone-B (the second drone source) errors are 0.15–0.4 m/s and, remarkably, *lower* at 3–4 m/s than when hovering; drone-A (racing drone) errors are 0.45–1.0 m/s at every speed.

Two data findings explain this, both measured on the released train/val files. (1) *Frame convention of drone-B*. Its gyroscope z axis is inverted relative to the quaternion/velocity body frame (ground-truth yaw rate = $-\omega_z$, $R^2=0.97$), and its accelerometer dynamics match the ground-truth specific force under no axis relabelling ($R^2<0$ at 2 Hz, unconstrained fits give axis gains of 0.6–2.1), whereas the other three platforms and drone-A have identity conventions with R^2 0.97–0.99. The network still learns drone-B well because the source is trivially identifiable from the signal (high-frequency accelerometer content differs by $10\times$), but the physics it learns there does not transfer to drone-A, which has only ≈ 30 min of data. (2) *Racing-drone recordings*. Every fast flight appears twice, recorded by two IMUs (one level, one pitched $\approx 45^\circ$, e.g. `drone_val.0005/0006`). Their IMU streams are offset from the ground truth by *different* amounts (≈ -10 ms for one IMU type, $-40 \dots -70$ ms for the other, measured by cross-correlating $|\omega|$ with the quaternion-derived rate), and after time alignment and the best-fit rotation the two recordings’ velocity targets still disagree by 0.15–0.9 m/s ($\approx 5^\circ$ of mount inconsistency at 4–7 m/s). This is a floor on that source that no IMU-only model can pass, and about a quarter of the test drone trajectories are of this type. Several dog trajectories also have offsets of 40–150 ms. We found no evidence of negative transfer between platforms: every platform improved with every improvement of the shared recipe, and removing a platform never helped another.




Fig. 3. Per-trajectory val AVE of the final recipe (train-only run) against mean speed. The racing-drone source is an outlier at every speed; the slowest drone-B sequences sit at the floor set by hover drift.

E. Failure-Mode Analysis

(i) Racing-drone flights: errors of 0.5–1 m/s that are bias-like per recording, consistent with the timing and mount inconsistencies above. (ii) Slow hover: velocity variations of ± 0.15 m/s at 0.1 Hz produce ≈ 0.1 m/s² of acceleration under 4–6 m/s² of vibration – unobservable. (iii) Fast cars (> 3 m/s, 25 windows in val) have errors of 3–4 m/s: this regime does not exist in the training data. Unresolved: a principled way to make the drone-B convention consistent with the rest of the data (it is not a fixed axis relabelling), and any use of the two-IMU redundancy.

V. INSIGHTS AND LESSONS LEARNED

A single universal model is feasible for the part of the problem that is observable. One 3.9 M-parameter network with no platform input reaches 0.07–0.13 m/s AVE on car, dog and human and ≈ 0.27 m/s on ordinary drone flight, and every change that helped one platform helped the others. The embodiment is inferred implicitly and reliably; specialization was never needed. **The bottleneck is data quality and observability, not the model.** Our largest remaining error comes from a source whose labels are internally inconsistent at the level of tens of milliseconds and a few degrees; more capacity, more epochs and more augmentation all left it untouched, while everything else kept improving. **Context transfers; conventions do not.** Reading whole trajectories, metric-shaped sampling and physically exact augmentation are sensor- and platform-agnostic and should transfer to other IMUs and vehicles; anything learned about a specific source’s axis convention or timing offset will not. **Per hour of data, slow-and-clean beats fast-and-noisy:** 30 min of racing-drone data contributed almost nothing to generalization, while the drone-B hours drove the drone-B error down steadily. **Evaluation.** With ~ 45 drone test trajectories averaged per trajectory, a public split is dominated by a few short racing flights; we measured ± 0.01 –0.02 leaderboard differences between identical recipes. Promising directions: physics-informed architectures that estimate bias, gravity and time offset explicitly with uncertainty; self-supervised consistency between the two IMUs of the same flight; and datasets with verified time synchronization.

A. Feedback to the Organizers

The rules, metric and starter kit were unusually clear, and the exact scorer in the kit made local validation trustworthy (our val floor 0.013 and all-zero 1.015 matched the published references). What cost us the most time was discovering the data issues of Sec. IV-D ourselves; documenting per-source conventions, IMU types and synchronization quality (or fixing the drone-B gyro sign and the racing-drone offsets) would let teams work on the model instead. The metric behaved as documented, but the per-trajectory averaging makes short racing-drone clips worth as much as 30-minute car drives, so the public leaderboard is noisy; averaging per window within a platform, or a larger public split, would reduce it. The second-IMU duplicates of the same flights should be either merged or declared. For the next edition, causal/online evaluation and a held-out platform would measure generalization more directly than a platform-blind test set whose sources are identifiable from the signal.

VI. CONCLUSION

A trajectory-context network trained with a metric-shaped sampler and physically exact augmentation gives a single unified inertial velocity estimator that is accurate on three of the four platforms and on ordinary drone flight, reaching 0.219 on the full test set (0.286 on the harder public split, from our own starting point of 0.553) with one weight set and no post-processing. The remaining error is dominated by

one drone source whose ground truth is internally inconsistent; closing it requires better data, or models that estimate synchronization and mounting explicitly.

ACKNOWLEDGMENT

We thank the CMU AirLab organizers for the dataset, the exact scorer and the scoring service, and Kaggle and Google Colab for compute. The AI assistant Claude (Anthropic) was used for code development and analysis under the authors’ direction.

APPENDIX

A. Per-Sequence Results (Full Table)

B. Reproducibility and Submission Artifacts

Kaggle team **Hack2Publish**. Ranked submission **56402554** (submission_final_kaggle_w_uni.csv, 2026-09-20 17:34 UTC, md5

fd12d4f25d7b52d44aaa4e4a14534378, public 0.28609). Second selected entry **56332266** (submission_final_v4.160.csv, md5 97a6e86dba7dfc19929e050bef5857b0, public 0.28696): the width-128 / 2-layer variant trained on an Apple M5 with length-weighted sampling. Code: the GitHub repository below (training commit a9a5384; the Kaggle kernel script that trained the ranked checkpoint is in kaggle.kernel/). Checkpoints: the Hugging Face repository below – model.pt (SHA-256 69764745d0589eba846d571b364aa3d3 3fe6a25be5e516879339b3231ceeb9b1) and model_secondary_final_v4.160.pt (f1eb2086...), with infer.py, model.py, pinned requirements.txt (torch 2.14.0, numpy 2.5.2, pandas 3.0.5) and SHA256SUMS; the exact argparse configuration is stored inside each checkpoint and in configs/. Seed 0; every reported score is a single run. Run-to-run spread: the identical recipe trained on two different GPUs scored 0.2870 and 0.3052 publicly, and the ranked recipe 0.2861 (Kaggle GPU) vs 0.2876 (laptop). Inference reproduces the submitted CSV to ≤ 0.005 m/s per window across devices (TF32 disabled).

Compliance declaration.

- 1) **Yes**. All predictions come from a single model with one shared set of weights.
- 2) **No** platform classification, routing or specialization at inference. An auxiliary platform head exists in the training graph only; its output is never computed at inference and nothing depends on it.
- 3) **No** model ensembling and **no** test-time augmentation. Weight averaging within a single run only: EMA of the weights during training and a uniform average of the EMA weights over the last 20 epochs.
- 4) Test data was read only by the inference script. All architecture and hyper-parameter choices were made on the released val split with the organizers’ scorer; the final model is a fixed-schedule refit on train+val with no checkpoint selection. Among the finished

refits, the two Kaggle entries were chosen by val-validated recipe first and public score second; no training or post-processing was tuned against the leaderboard.

- 5) The development-time system and the submitted system are identical. The only differences between the ranked and the second entry are the trunk width / transformer depth and the sampling weights, as stated above.

C. Final Predictions

The ranked CSV (30,644 rows, columns window_id, vx, vy, vz, md5 above) is attached to the Submission Form and is regenerated bit-exactly on the training device from the Hugging Face repository by `python infer.py --traj_dir test --windows index/test_windows.csv --out submission.csv`. No post-processing is applied after the network output beyond the overlap averaging described in Sec. II-E; there are no missing or degenerate windows.

D. Model and Inference Code

Weights, inference entry point, configuration and environment (Hugging Face):

<https://huggingface.co/mdhamidhosen/tartanimu-unified-hosen42>

Training code, analysis scripts and this report’s sources (GitHub):

<https://github.com/hamidhosen42/TartanIMU-Challenge-Multi-Platform-Inertial-Odometry>

REFERENCES

- [1] S. Zhao, S. Zhou, R. Blanchard, Y. Qiu, W. Wang, and S. Scherer, “Tartan IMU: A Light Foundation Model for Inertial Positioning in Robotics,” in *Proc. IEEE/CVF CVPR*, 2025, pp. 22520–22529.
- [2] J. Sturm, N. Engelhard, F. Endres, W. Burgard, and D. Cremers, “A Benchmark for the Evaluation of RGB-D SLAM Systems,” in *Proc. IEEE/RSJ IROS*, 2012, pp. 573–580.
- [3] S. Herath, H. Yan, and Y. Furukawa, “RoNIN: Robust Neural Inertial Navigation in the Wild,” in *Proc. IEEE ICRA*, 2020.
- [4] W. Liu, D. Caruso, E. Ilg, J. Dong, A. Angelova, and S. Belongie, “TLIO: Tight Learned Inertial Odometry,” *IEEE RA-L*, vol. 5, no. 4, 2020.

TABLE V

PER-SEQUENCE CHALLENGE RESULTS OVER THE 89 TEST SEQUENCES (OFFICIAL PER-SEQUENCE SCORING SERVICE; SUBMISSION MD5 FD12D4F25D7B52D44AAA4E4A14534378). **RTE** (RELATIVE TRAJECTORY ERROR, M) FOLLOWS STURM *et al.* [2]: OVER SLIDING WINDOWS OF $\Delta t = 5$ s, THE RMSE OF THE DIFFERENCE BETWEEN PREDICTED AND GROUND-TRUTH RELATIVE DISPLACEMENT,
$$\text{RTE} = \sqrt{\frac{1}{n} \sum_i \|(\hat{\mathbf{p}}_{i+\Delta t} - \hat{\mathbf{p}}_i) - (\mathbf{p}_{i+\Delta t} - \mathbf{p}_i)\|^2}$$
. SEQUENCE IDS FOLLOW THE SCORING SERVICE'S ORDERING.

Seq.	ATE ₂₀	AVE	RTE	Seq.	ATE ₂₀	AVE	RTE	Seq.	ATE ₂₀	AVE	RTE
1	0.830	0.337	1.873	31	0.703	0.298	1.605	61	0.906	0.360	1.158
2	0.498	0.351	1.181	32	0.319	0.084	0.246	62	0.073	0.020	0.061
3	0.535	0.205	0.765	33	0.210	0.051	0.207	63	0.455	0.070	0.327
4	0.498	0.308	0.903	34	1.581	0.366	1.800	64	0.404	0.059	0.268
5	0.905	0.405	1.835	35	1.056	0.413	1.659	65	0.106	0.038	0.130
6	0.674	0.089	0.763	36	1.016	0.325	1.878	66	2.153	0.226	1.577
7	1.116	0.344	1.508	37	2.555	2.108	11.718	67	0.113	0.032	0.143
8	0.468	0.039	0.208	38	2.152	1.807	8.917	68	0.526	0.048	0.222
9	0.640	0.278	0.972	39	0.644	0.256	0.977	69	0.686	0.084	0.440
10	0.771	0.086	0.693	40	0.784	0.318	1.853	70	0.503	0.073	0.330
11	1.024	0.658	2.017	41	1.081	0.327	2.140	71	0.196	0.060	0.192
12	0.307	0.154	0.672	42	0.493	0.088	0.483	72	0.851	0.452	1.437
13	0.300	0.044	0.262	43	0.920	0.286	1.115	73	0.807	0.131	0.767
14	0.564	0.383	1.228	44	0.212	0.051	0.315	74	0.595	0.310	1.006
15	1.021	0.073	0.409	45	0.083	0.053	0.126	75	0.548	0.393	1.216
16	0.127	0.034	0.224	46	0.258	0.059	0.241	76	0.399	0.097	0.423
17	1.136	0.559	2.669	47	0.840	0.387	1.063	77	1.166	0.567	2.337
18	1.360	0.565	2.554	48	0.211	0.015	0.101	78	0.197	0.029	0.098
19	0.521	0.294	1.308	49	0.798	0.255	1.343	79	0.817	0.261	1.419
20	0.140	0.053	0.370	50	0.669	0.210	1.117	80	0.448	0.174	0.910
21	0.706	0.209	1.133	51	0.080	0.037	0.097	81	0.645	0.116	0.474
22	0.573	0.391	1.197	52	0.258	0.070	0.359	82	2.896	0.389	2.533
23	1.027	0.241	1.984	53	0.394	0.279	0.513	83	0.449	0.074	0.372
24	0.392	0.043	0.193	54	0.222	0.070	0.260	84	0.477	0.057	0.321
25	0.105	0.029	0.109	55	0.269	0.144	0.665	85	0.074	0.032	0.097
26	0.596	0.086	0.353	56	3.406	2.896	15.349	86	0.305	0.027	0.148
27	0.476	0.060	0.290	57	0.922	0.356	2.145	87	0.414	0.173	0.657
28	0.566	0.256	1.271	58	0.219	0.050	0.250	88	1.028	0.294	2.012
29	3.661	2.886	16.039	59	0.361	0.074	0.385	89	0.148	0.028	0.091
30	1.136	0.548	2.608	60	0.156	0.066	0.214				